

[link text](#)D.LAKSHMAN 21BCE9053 CSE(AI/ML)

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
data=pd.read_csv("Employee-Attrition.csv")
```

```
data.head()
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	Empl
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	1	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medical	1	

5 rows × 35 columns

```
data.tail()
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	Empl
1465	36	No	Travel_Frequently	884	Research & Development	23	2	Medical	1	
1466	39	No	Travel_Rarely	613	Research & Development	6	1	Medical	1	
1467	27	No	Travel_Rarely	155	Research & Development	4	3	Life Sciences	1	
1468	49	No	Travel_Frequently	1023	Sales	2	3	Medical	1	
1469	34	No	Travel_Rarely	628	Research & Development	8	3	Medical	1	

5 rows × 35 columns

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   Age              1470 non-null    int64  
 1   Attrition        1470 non-null    object 
 2   BusinessTravel   1470 non-null    object 
 3   DailyRate        1470 non-null    int64  
 4   Department       1470 non-null    object 
 5   DistanceFromHome 1470 non-null    int64  
 6   Education        1470 non-null    int64  
 7   EducationField   1470 non-null    object 
 8   EmployeeCount    1470 non-null    int64  
 9   EmployeeNumber   1470 non-null    int64  
 10  EnvironmentSatisfaction 1470 non-null    int64  
 11  Gender            1470 non-null    object 
 12  HourlyRate       1470 non-null    int64  
 13  JobInvolvement   1470 non-null    int64  
 14  JobLevel          1470 non-null    int64  
 15  JobRole           1470 non-null    object 
 16  JobSatisfaction  1470 non-null    int64  
 17  MaritalStatus     1470 non-null    object 
 18  MonthlyIncome     1470 non-null    int64  
 19  MonthlyRate       1470 non-null    int64
```

```

20 NumCompaniesWorked      1470 non-null   int64
21 Over18                  1470 non-null   object
22 Overtime                 1470 non-null   object
23 PercentSalaryHike        1470 non-null   int64
24 PerformanceRating         1470 non-null   int64
25 RelationshipSatisfaction 1470 non-null   int64
26 StandardHours             1470 non-null   int64
27 StockOptionLevel          1470 non-null   int64
28 TotalWorkingYears          1470 non-null   int64
29 TrainingTimesLastYear     1470 non-null   int64
30 WorkLifeBalance            1470 non-null   int64
31 YearsAtCompany             1470 non-null   int64
32 YearsInCurrentRole        1470 non-null   int64
33 YearsSinceLastPromotion    1470 non-null   int64
34 YearsWithCurrManager       1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB

```

```
data.describe()
```

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	EnvironmentSatisfaction	Hour
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470.000000	1470.000000	1470.
mean	36.923810	802.485714	9.192517	2.912925	1.0	1024.865306	2.721769	65.
std	9.135373	403.509100	8.106864	1.024165	0.0	602.024335	1.093082	20.
min	18.000000	102.000000	1.000000	1.000000	1.0	1.000000	1.000000	30.
25%	30.000000	465.000000	2.000000	2.000000	1.0	491.250000	2.000000	48.
50%	36.000000	802.000000	7.000000	3.000000	1.0	1020.500000	3.000000	66.
75%	43.000000	1157.000000	14.000000	4.000000	1.0	1555.750000	4.000000	83.
max	60.000000	1499.000000	29.000000	5.000000	1.0	2068.000000	4.000000	100.

8 rows × 26 columns

▼ HANDLING NULL VALUES

```
data.isnull().any()
```

```

Age                False
Attrition          False
BusinessTravel     False
DailyRate           False
Department          False
DistanceFromHome   False
Education           False
EducationField      False
EmployeeCount       False
EmployeeNumber      False
EnvironmentSatisfaction False
Gender              False
HourlyRate          False
JobInvolvement      False
JobLevel             False
JobRole              False
JobSatisfaction     False
MaritalStatus        False
MonthlyIncome        False
MonthlyRate          False
NumCompaniesWorked  False
Over18               False
Overtime             False
PercentSalaryHike    False
PerformanceRating    False
RelationshipSatisfaction False
StandardHours        False
StockOptionLevel     False
TotalWorkingYears    False
TrainingTimesLastYear False
WorkLifeBalance      False
YearsAtCompany        False
YearsInCurrentRole    False
YearsSinceLastPromotion False
YearsWithCurrManager False
dtype: bool

```

```
data.isnull().sum()

Age          0
Attrition    0
BusinessTravel 0
DailyRate     0
Department    0
DistanceFromHome 0
Education     0
EducationField 0
EmployeeCount 0
EmployeeNumber 0
EnvironmentSatisfaction 0
Gender        0
HourlyRate    0
JobInvolvement 0
JobLevel      0
JobRole       0
JobSatisfaction 0
MaritalStatus 0
MonthlyIncome 0
MonthlyRate    0
NumCompaniesWorked 0
Over18        0
OverTime      0
PercentSalaryHike 0
PerformanceRating 0
RelationshipSatisfaction 0
StandardHours 0
StockOptionLevel 0
TotalWorkingYears 0
TrainingTimesLastYear 0
WorkLifeBalance 0
YearsAtCompany 0
YearsInCurrentRole 0
YearsSinceLastPromotion 0
YearsWithCurrManager 0
dtype: int64
```

```
cor=data.corr()
```

```
C:\Users\srich\AppData\Local\Temp\ipykernel_26344\1426905697.py:1: FutureWarning: The default value of numeric_only in DataFrame.corr is
```

```
cor=data.corr()
```

```
fig=plt.figure(figsize=(18,18))
sns.heatmap(cor,annot=True)
```

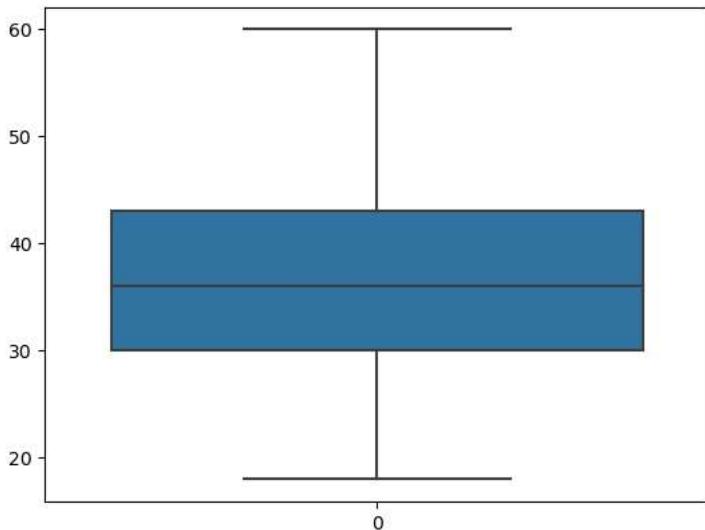
<Axes: >



▼ OUTLIERS

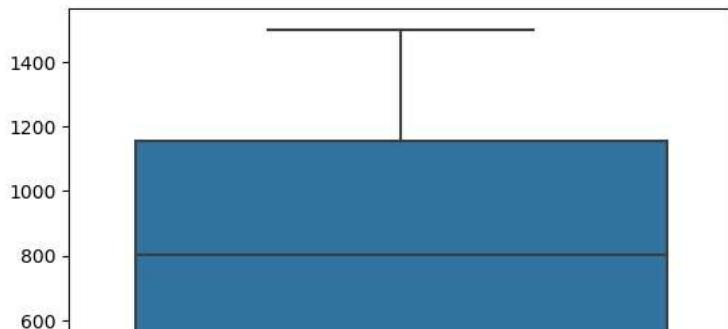
sns.boxplot(data["Age"])

<Axes: >



sns.boxplot(data["DailyRate"])

<Axes: >



data.describe()

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	EnvironmentSatisfaction	Hour
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470.000000	1470.000000	1470.
mean	36.923810	802.485714	9.192517	2.912925	1.0	1024.865306	2.721769	65.
std	9.135373	403.509100	8.106864	1.024165	0.0	602.024335	1.093082	20.
min	18.000000	102.000000	1.000000	1.000000	1.0	1.000000	1.000000	30.
25%	30.000000	465.000000	2.000000	2.000000	1.0	491.250000	2.000000	48.
50%	36.000000	802.000000	7.000000	3.000000	1.0	1020.500000	3.000000	66.
75%	43.000000	1157.000000	14.000000	4.000000	1.0	1555.750000	4.000000	83.
max	60.000000	1499.000000	29.000000	5.000000	1.0	2068.000000	4.000000	100.

8 rows × 26 columns

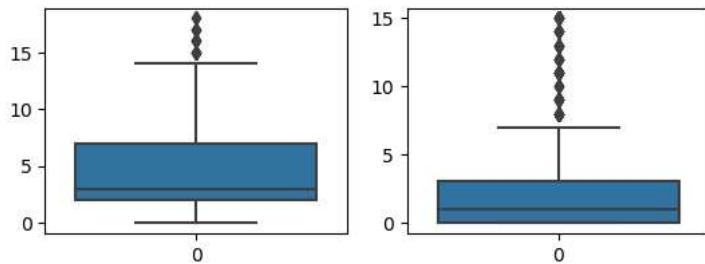
data.head()

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EducationField	EmployeeCount	Empl
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences	1	
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences	1	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Other	1	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences	1	
4	27	No	Travel_Rarely	591	Research & Development	2	1	Medical	1	

5 rows × 35 columns

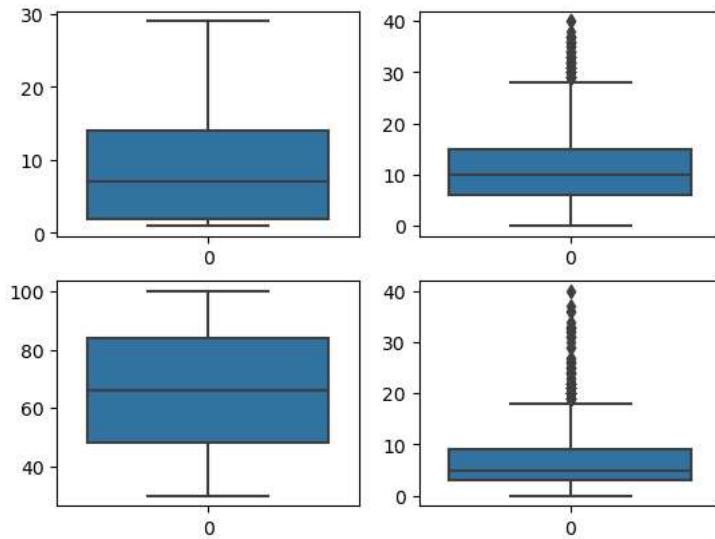
```
fig, axes = plt.subplots(2,2)
sns.boxplot(data=data["YearsInCurrentRole"],ax=axes[0,0])
sns.boxplot(data=data["YearsSinceLastPromotion"],ax=axes[0,1])
sns.boxplot(data=data["YearsWithCurrManager"],ax=axes[1,0])
sns.boxplot(data=data["WorkLifeBalance"],ax=axes[1,1])
```

<Axes: >



```
fig, axes = plt.subplots(2,2)
sns.boxplot(data=data["DistanceFromHome"],ax=axes[0,0])
sns.boxplot(data=data["TotalWorkingYears"],ax=axes[0,1])
sns.boxplot(data=data["HourlyRate"],ax=axes[1,0])
sns.boxplot(data=data["YearsAtCompany"],ax=axes[1,1])
```

<Axes: >



▼ HANDLING THE OUTLIERS

```
YearsInCurrentRole_q1 = data.YearsInCurrentRole.quantile(0.25)
YearsInCurrentRole_q3 = data.YearsInCurrentRole.quantile(0.75)
IQR_YearsInCurrentRole=YearsInCurrentRole_q3-YearsInCurrentRole_q1
upperlimit_YearsInCurrentRole=YearsInCurrentRole_q3+1.5*IQR_YearsInCurrentRole
lower_limit_YearsInCurrentRole =YearsInCurrentRole_q1-1.5*IQR_YearsInCurrentRole
median_YearsInCurrentRole=data["YearsInCurrentRole"].median()
data['YearsInCurrentRole'] = np.where(
    (data['YearsInCurrentRole'] > upperlimit_YearsInCurrentRole),
    median_YearsInCurrentRole,
    data['YearsInCurrentRole']
)

YearsSinceLastPromotion_q1 = data.YearsSinceLastPromotion.quantile(0.25)
YearsSinceLastPromotion_q3 = data.YearsSinceLastPromotion.quantile(0.75)
IQR_YearsSinceLastPromotion=YearsSinceLastPromotion_q3-YearsSinceLastPromotion_q1
upperlimit_YearsSinceLastPromotion=YearsSinceLastPromotion_q3+1.5*IQR_YearsSinceLastPromotion
lower_limit_YearsSinceLastPromotion =YearsSinceLastPromotion_q1-1.5*IQR_YearsSinceLastPromotion
median_YearsSinceLastPromotion=data["YearsSinceLastPromotion"].median()
data['YearsSinceLastPromotion'] = np.where(
    (data['YearsSinceLastPromotion'] > upperlimit_YearsSinceLastPromotion),
    median_YearsSinceLastPromotion,
    data['YearsSinceLastPromotion']
)

YearsWithCurrManager_q1 = data.YearsWithCurrManager.quantile(0.25)
YearsWithCurrManager_q3 = data.YearsWithCurrManager.quantile(0.75)
IQR_YearsWithCurrManager=YearsWithCurrManager_q3-YearsWithCurrManager_q1
upperlimit_YearsWithCurrManager=YearsWithCurrManager_q3+1.5*IQR_YearsWithCurrManager
```

```

lower_limit_YearsWithCurrManager = YearsWithCurrManager_q1 - 1.5 * IQR_YearsWithCurrManager
median_YearsWithCurrManager = data["YearsWithCurrManager"].median()
data['YearsWithCurrManager'] = np.where(
    (data['YearsWithCurrManager'] > upperlimit_YearsWithCurrManager),
    median_YearsWithCurrManager,
    data['YearsWithCurrManager']
)

TotalWorkingYears_q1 = data.TotalWorkingYears.quantile(0.25)
TotalWorkingYears_q3 = data.TotalWorkingYears.quantile(0.75)
IQR_TotalWorkingYears = TotalWorkingYears_q3 - TotalWorkingYears_q1
upperlimit_TotalWorkingYears = TotalWorkingYears_q3 + 1.5 * IQR_TotalWorkingYears
lower_limit_TotalWorkingYears = TotalWorkingYears_q1 - 1.5 * IQR_TotalWorkingYears
median_TotalWorkingYears = data["TotalWorkingYears"].median()
data['TotalWorkingYears'] = np.where(
    (data['TotalWorkingYears'] > upperlimit_TotalWorkingYears),
    median_TotalWorkingYears,
    data['TotalWorkingYears']
)

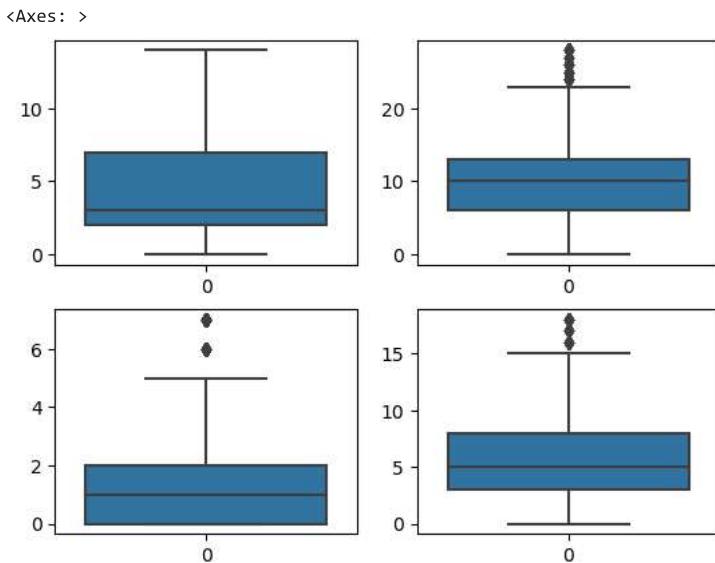
YearsAtCompany_q1 = data.YearsAtCompany.quantile(0.25)
YearsAtCompany_q3 = data.YearsAtCompany.quantile(0.75)
IQR_YearsAtCompany = YearsAtCompany_q3 - YearsAtCompany_q1
upperlimit_YearsAtCompany = YearsAtCompany_q3 + 1.5 * IQR_YearsAtCompany
lower_limit_YearsAtCompany = YearsAtCompany_q1 - 1.5 * IQR_YearsAtCompany
median_YearsAtCompany = data["YearsAtCompany"].median()
data['YearsAtCompany'] = np.where(
    (data['YearsAtCompany'] > upperlimit_YearsAtCompany),
    median_YearsAtCompany,
    data['YearsAtCompany']
)

```

```

fig, axes = plt.subplots(2,2)
sns.boxplot(data=data["YearsWithCurrManager"], ax=axes[0,0])
sns.boxplot(data=data["TotalWorkingYears"], ax=axes[0,1])
sns.boxplot(data=data["YearsSinceLastPromotion"], ax=axes[1,0])
sns.boxplot(data=data["YearsAtCompany"], ax=axes[1,1])

```



```
data.head()
```

```
Age Attrition BusinessTravel DailyRate Department DistanceFromHome Education EducationField EmployeeCount Empl
0 41 Yes Travel_Rarely 1102 Sales 1 2 Life Sciences 1
1 49 No Travel_Frequently 279 Research & Development 8 1 Life Sciences 1
Research &
data.drop("EducationField",axis=1,inplace=True)
Research &
data.head(2)

Age Attrition BusinessTravel DailyRate Department DistanceFromHome Education EmployeeCount EmployeeNumber Envir
0 41 Yes Travel_Rarely 1102 Sales 1 2 1 1
1 49 No Travel_Frequently 279 Research & Development 8 1 1 2
2 rows x 34 columns
```

```
array(['Travel_Rarely', 'Travel_Frequently', 'Non-Travel'], dtype=object)
```

▼ SPLITTING THE DATA

```
y=data["Attrition"]  
y.head()  
  
0    Yes  
1    No  
2    Yes  
3    No  
4    No  
Name: Attrition
```

```
data.drop("Attrition",axis=1,inplace=True)
```

```
data.head()
```

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EmployeeCount	EmployeeNumber	EnvironmentSati
0	41	Travel_Rarely	1102	Sales	1	2	1	1	
1	49	Travel_Frequently	279	Research & Development	8	1	1	2	
2	37	Travel_Rarely	1373	Research & Development	2	2	1	4	
3	33	Travel_Frequently	1392	Research & Development	3	4	1	5	
4	27	Travel_Rarely	591	Research & Development	2	1	1	7	

5 rows × 33 columns

▼ ENCODING

```
from sklearn.preprocessing import LabelEncoder  
  
le=LabelEncoder()  
  
data[\"BusinessTravel\"] = le.fit_transform(data[\"BusinessTravel\"])
```

```

data["Department"] = le.fit_transform(data["Department"])

data["Gender"] = le.fit_transform(data["Gender"])

y = le.fit_transform(y)

array([1, 0, 1, ..., 0, 0, 0])

data["JobRole"] = le.fit_transform(data["JobRole"])

data["Over18"] = le.fit_transform(data["Over18"])

data["MaritalStatus"] = le.fit_transform(data["MaritalStatus"])

data["OverTime"] = le.fit_transform(data["OverTime"])

data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 33 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Age              1470 non-null    int64  
 1   BusinessTravel   1470 non-null    int32  
 2   DailyRate        1470 non-null    int64  
 3   Department       1470 non-null    int32  
 4   DistanceFromHome 1470 non-null    int64  
 5   Education        1470 non-null    int64  
 6   EmployeeCount    1470 non-null    int64  
 7   EmployeeNumber   1470 non-null    int64  
 8   EnvironmentSatisfaction 1470 non-null    int64  
 9   Gender            1470 non-null    int32  
 10  HourlyRate       1470 non-null    int64  
 11  JobInvolvement   1470 non-null    int64  
 12  JobLevel          1470 non-null    int64  
 13  JobRole           1470 non-null    int32  
 14  JobSatisfaction  1470 non-null    int64  
 15  MaritalStatus    1470 non-null    int32  
 16  MonthlyIncome    1470 non-null    int64  
 17  MonthlyRate      1470 non-null    int64  
 18  NumCompaniesWorked 1470 non-null    int64  
 19  Over18           1470 non-null    int32  
 20  OverTime          1470 non-null    int32  
 21  PercentSalaryHike 1470 non-null    int64  
 22  PerformanceRating 1470 non-null    int64  
 23  RelationshipSatisfaction 1470 non-null    int64  
 24  StandardHours    1470 non-null    int64  
 25  StockOptionLevel 1470 non-null    int64  
 26  TotalWorkingYears 1470 non-null    float64 
 27  TrainingTimesLastYear 1470 non-null    int64  
 28  WorkLifeBalance   1470 non-null    int64  
 29  YearsAtCompany    1470 non-null    float64 
 30  YearsInCurrentRole 1470 non-null    float64 
 31  YearsSinceLastPromotion 1470 non-null    float64 
 32  YearsWithCurrManager 1470 non-null    float64 
dtypes: float64(5), int32(7), int64(21)
memory usage: 338.9 KB

```

▼ TRAIN TEST SPLIT

```

from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(data, y, test_size=0.3, random_state=0)

x_train.shape, x_test.shape, y_train.shape, y_test.shape

((1029, 33), (441, 33), (1029,), (441,))

```

▼ FEATURE SCALING

```
from sklearn.preprocessing import StandardScaler

sc=StandardScaler()

x_train=sc.fit_transform(x_train)

x_test=sc.fit_transform(x_test)
```

BUILDING THE MODEL

▼ MULTI LINEAR REGRESSION

```
from sklearn.linear_model import LinearRegression

lr = LinearRegression()

lr.fit(x_train,y_train)



▶ LinearRegression  

    LinearRegression()



lr.coef_ #slope(m)

array([-3.54982205e-02,  5.83302672e-05, -1.72249253e-02,  3.46305828e-02,
       2.45263256e-02,  3.89940919e-03, -8.89214800e+11, -9.43596046e-03,
      -4.12143211e-02,  1.05731111e-02, -3.10438176e-03, -3.84488534e-02,
     -1.53491202e-02, -1.57440821e-02, -3.66214700e-02,  3.35650229e-02,
     -5.85270940e-03,  5.89670852e-03,  3.77895930e-02, -1.09571684e+07,
      9.55711035e-02, -2.54155836e-02,  1.99397870e-02, -2.64643443e-02,
     -2.64138305e+04, -1.79073382e-02, -3.31099987e-02, -1.08374986e-02,
     -3.09087610e-02, -2.49869759e-02, -1.08495820e-02,  2.10399164e-02,
     -6.46780261e-03])
```

```
lr.intercept_ #(c)

0.16229348882410102
```

```
y_pred = lr.predict(x_test)
```

```
y_pred

array([ 1.30796232e-01,  2.17900848e-01,  3.46642078e-01,  5.58365807e-03,
       4.99016900e-01,  1.01768656e-01,  3.45454856e-01,  1.23487843e-01,
      -1.60508211e-01,  4.02774192e-01,  1.43863626e-01,  2.67493211e-01,
     -4.61376545e-02,  5.58563347e-01,  2.81985888e-01,  1.48533683e-02,
     1.78513543e-01,  2.78218063e-01,  9.39501521e-02,  2.16824147e-01,
     2.65686758e-01,  1.40480328e-02,  8.35045241e-02,  9.65278467e-02,
     5.10152592e-01,  2.94947412e-01,  7.87474989e-02,  1.26281046e-01,
     5.05599470e-01,  8.44181518e-02, -7.94455244e-02,  2.00664655e-02,
     1.07159698e-01,  3.65801296e-01,  1.25083285e-01,  5.15443890e-02,
     1.06691263e-01,  6.10560784e-02,  6.60004307e-02,  4.86894584e-02,
     -1.07548569e-02, -2.94781175e-02,  5.20898060e-02, -1.58276250e-02,
     -1.77848612e-02,  4.18379396e-01,  3.66954632e-01, -2.14407437e-01,
     5.48302265e-01,  4.40485603e-01,  1.97011357e-01,  4.42145584e-01,
     1.46307401e-01,  3.75270817e-01,  4.93018248e-01,  2.95973962e-01,
     -4.67912123e-02,  3.16682815e-01, -7.70189517e-03,  2.53080649e-01,
     -3.14434764e-02,  2.83109378e-01,  9.08930998e-02,  1.26210137e-01,
     3.60059433e-01,  2.42994338e-02,  3.55930832e-01,  1.96149466e-01,
     1.27766995e-01,  1.19157288e-01, -2.87607873e-02,  3.18174064e-01,
     1.08182581e-01,  1.25900358e-01,  2.30252390e-01,  9.79248506e-02,
     9.13876026e-02,  2.72922895e-01,  2.52302174e-01,  4.06448037e-02,
     -9.11781684e-02, -1.11336412e-02,  1.94492552e-01, -2.23451867e-02,
     -1.72239428e-02,  1.16255875e-01,  8.35014963e-02,  2.28958407e-03,
     4.87956812e-02,  2.40960805e-01,  3.14231768e-01,  2.25370717e-01,
     3.30101918e-01,  2.38617684e-01, -2.14125086e-02,  2.27066280e-01,
     3.01683445e-01,  2.98737390e-01,  9.82670241e-02,  8.83243788e-02,
     2.86155158e-01,  4.99817815e-01,  3.04309363e-01, -5.27010704e-03,
     1.71351185e-01, -5.95768064e-03,  2.56475095e-02,  2.15399623e-01,
     6.10705957e-02,  1.64259248e-01,  1.08747531e-01,  1.07853987e-01,
```

3.13741410e-02,	1.96152859e-01,	9.68496302e-02,	3.33338078e-02,
1.06769816e-01,	2.33826858e-01,	-8.62611908e-02,	-7.73660020e-02,
2.00264727e-01,	3.42215765e-02,	1.28025570e-01,	6.03549521e-01,
5.90092635e-03,	-3.09303352e-02,	-1.46523580e-01,	2.18857517e-01,
2.75684912e-01,	1.68178132e-01,	-2.73302560e-03,	6.21705896e-01,
4.40410881e-01,	3.95659751e-01,	1.69391826e-01,	4.18498413e-01,
4.90851141e-01,	2.02156908e-01,	1.58125659e-01,	3.60608308e-01,
2.26116966e-01,	1.46080493e-01,	2.13446831e-01,	2.66808395e-01,
3.11988347e-01,	-8.15460457e-04,	1.49618075e-01,	-1.34986580e-01,
2.07992802e-01,	2.79499979e-01,	1.16456567e-01,	2.74369274e-01,
5.51722773e-02,	3.41682216e-01,	1.70237398e-01,	-3.43112640e-04,
-4.05339167e-02,	1.34919977e-01,	-1.02695835e-01,	-5.64532727e-02,
3.37155137e-01,	-9.52131633e-02,	2.11183014e-01,	6.18254327e-01,
2.03804707e-01,	3.04259774e-01,	1.82073591e-01,	1.84394143e-01,
-2.87766822e-03,	-8.96366447e-02,	4.16941567e-02,	1.30414229e-01,
1.73722893e-01,	1.58172468e-01,	-8.67312423e-02,	1.87531180e-01,
1.99848543e-01,	1.82435808e-01,	1.03617628e-01,	1.91507185e-01,
2.51148859e-03,	1.95023436e-01,	-6.13160664e-02,	5.85144923e-01,
6.61570629e-02,	4.50721221e-02,	3.30228796e-01,	9.79202044e-02,
5.51504673e-01,	1.52340046e-01,	3.58897737e-01,	3.66200721e-01,
2.47443494e-01,	5.94731854e-02,	1.28084751e-01,	2.80510398e-01,
7.23178238e-02,	-8.08554776e-02,	3.39348594e-01,	8.20383069e-02,
2.19295711e-01,	2.48497484e-01,	4.968380400e-01,	1.36684877e-01,
2.88225529e-01,	4.60162124e-02,	4.52391987e-01,	-8.24840394e-02,
2.26359015e-01,	1.41222376e-02,	1.61414983e-01,	2.32302068e-01,
9.12373829e-02,	1.19076584e-01,	2.12036068e-01,	-2.68369095e-02,
4.51730340e-02,	1.10311902e-01,	2.54707406e-02,	2.31690547e-01,
1.63046164e-01,	2.42643454e-01,	5.44862264e-01,	1.26147728e-01,
3.69618044e-01,	-8.09176638e-02,	1.41896523e-01,	2.86097884e-01,
3.14623385e-01,	1.32523750e-02,	-3.54222514e-02,	3.52906106e-01,

y_test

▼ LOGISTIC REGRESSION

```
from sklearn.linear_model import LogisticRegression
```

```
lg=LogisticRegression()
```

```
lg.fit(x_train,y_train)
```

```
‐ LogisticRegression  
LogisticRegression()
```

```
y_pred_lg=lg.predict(x_test)
```

y_pred

```
array([ 1.30796232e-01,  2.17900848e-01,  3.46642078e-01,  5.58365807e-03,
       4.99016900e-01,  1.01768656e-01,  3.45454856e-01,  1.23487843e-01,
      -1.60508211e-01,  4.02774192e-01,  1.43863626e-01,  2.67493211e-01,
      -4.61376545e-02,  5.58563347e-01,  2.81985888e-01,  1.48533683e-02,
       1.78513543e-01,  2.78218063e-01,  9.39515212e-02,  2.16823471e-01,
```

6.65686758e-01,	1.40480328e-02,	8.35045241e-02,	9.65278467e-02,
5.10152592e-01,	2.94947412e-01,	7.87474989e-02,	1.26281046e-01,
5.0559470e-01,	8.44181518e-02,	-7.94455244e-02,	2.00664655e-02,
1.07159698e-01,	3.65801296e-01,	1.25083285e-01,	5.15443890e-02,
1.06691263e-01,	6.18560784e-02,	6.66004307e-02,	4.86894584e-02,
-1.07548569e-02,	-2.94781175e-02,	5.20898060e-02,	-1.58276250e-02,
-1.77848612e-02,	4.18379396e-01,	3.66954632e-01,	-2.14407437e-01,
5.48302265e-01,	4.40485603e-01,	1.97011357e-01,	4.42145584e-01,
1.46307401e-01,	3.75270817e-01,	4.93018248e-01,	2.95973962e-01,
-4.67912123e-02,	3.16682815e-01,	-7.70189517e-03,	5.30808649e-01,
-3.14434764e-02,	2.83109378e-01,	9.08930998e-02,	1.26210137e-01,
3.60059433e-01,	2.42994338e-02,	3.55930832e-01,	1.96149466e-01,
1.27766995e-01,	1.19157288e-01,	-2.87607873e-02,	3.18174064e-01,
1.08182581e-01,	1.25900358e-01,	2.30252390e-01,	9.79248506e-02,
9.13876026e-02,	2.72922895e-01,	2.52302174e-01,	4.06448037e-02,
-9.11781684e-02,	-1.11336412e-02,	1.94492552e-01,	-2.23451867e-02,
-1.72239428e-02,	1.16255875e-01,	8.35014963e-02,	2.28958407e-03,
4.87956812e-02,	2.40960805e-01,	3.14231768e-01,	2.25370717e-01,
3.30101918e-01,	2.38617684e-01,	-2.14125086e-02,	2.27066280e-01,
3.01683445e-01,	2.98737390e-01,	9.82670241e-02,	8.83243788e-02,
2.86155158e-01,	4.99817815e-01,	3.04309363e-01,	-5.27010704e-03,
1.71351185e-01,	-5.95768064e-03,	2.56475095e-02,	1.25399623e-01,
6.10705957e-02,	1.64259248e-01,	1.08747531e-01,	1.07853987e-01,
-3.13741410e-02,	1.96152859e-01,	9.68496302e-02,	3.33338078e-02,
1.06769816e-01,	2.33826858e-01,	-8.62611908e-02,	-7.73660020e-02,
2.00264727e-01,	3.42215765e-02,	1.28025570e-01,	6.03549521e-01,
5.900926353e-03,	-3.09303352e-02,	-1.465235580e-01,	2.18857517e-01,
2.75684912e-01,	1.68178132e-01,	-2.73302560e-03,	2.61705896e-01,
4.40410881e-01,	3.95659751e-01,	1.69391826e-01,	4.18498413e-01,
4.90851141e-01,	2.02156908e-01,	1.58125659e-01,	3.60680308e-01,
2.26116966e-01,	1.46080493e-01,	2.13446831e-01,	2.66808395e-01,
3.11988347e-01,	-8.15460457e-04,	1.49618075e-01,	-1.34986507e-01,
2.07992802e-01,	2.79499979e-01,	1.16456567e-01,	2.74369274e-01,
5.51722773e-02,	3.41682216e-01,	1.70237398e-01,	-3.43112640e-04,
-4.05339167e-02,	1.34919977e-01,	-1.02695835e-01,	-5.64532727e-02,
3.37155137e-01,	-9.52131633e-02,	2.11183014e-01,	6.18254327e-01,
2.03804707e-01,	3.04259774e-01,	1.82073591e-01,	1.84394143e-01,
-2.87766822e-03,	-8.96366447e-02,	4.16941567e-02,	1.30414229e-01,
1.73722893e-01,	1.58172468e-01,	-8.67312423e-02,	1.87531180e-01,
1.99848543e-01,	1.82435808e-01,	1.03617628e-01,	1.91507185e-01,
2.51148859e-03,	1.95023436e-01,	-6.13160664e-02,	5.85144923e-01,
6.61570629e-02,	4.50721221e-02,	3.30228796e-01,	9.79202044e-02,
5.51504673e-01,	1.52340046e-01,	3.58897737e-01,	3.66200721e-01,
2.47443494e-01,	5.94731854e-02,	1.28084751e-01,	2.80510398e-01,
7.23178238e-02,	-8.08554776e-02,	3.39348594e-01,	8.20383069e-02,
2.19295711e-01,	2.48497484e-01,	4.96830400e-01,	1.36684877e-01,
2.88225529e-01,	4.601612124e-02,	4.52391987e-01,	-8.24840394e-02,
2.26359015e-01,	1.41222376e-02,	1.61414983e-01,	2.32302068e-01,
9.12373829e-02,	1.19076584e-01,	2.12036068e-01,	-2.68369095e-02,
4.51730340e-02,	1.10311902e-01,	2.54707406e-02,	2.31690547e-01,
1.63046164e-01,	2.42643454e-01,	5.44826624e-01,	1.26147728e-01,
3.69618044e-01,	-8.09176638e-02,	1.41896523e-01,	2.86097884e-01,
-2.14022326e-01,	1.23522757e-02,	2.54222511e-01,	2.50266162e-01,

y_test

```
score = lg.score(x_test, y_test)
print(score)
```

0.8820861678004536

▼ CONFUSION MATRIX

```
from sklearn import metrics
cm = metrics.confusion_matrix(y_test,y_pred_lg)
print(cm)

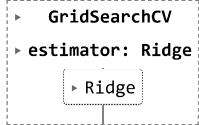
[[366  5]
 [ 47 23]]
```

▼ RIDGE AND LASSO

```
from sklearn.linear_model import Ridge
from sklearn.model_selection import GridSearchCV

rg=Ridge()

parametres={"alpha":[1,2,3,5,10,20,30,40,60,70,80,90]}
ridgecv=GridSearchCV(rg,parametres,scoring="neg_mean_squared_error",cv=5)
ridgecv.fit(x_train,y_train)
```



```
print(ridgecv.best_params_)

{'alpha': 90}

print(ridgecv.best_score_)

-0.1139062113923418

y_pred_rg=ridgecv.predict(x_test)

y_pred_rg
```

1.4946145e-02,	2.422843e-01,	-.2250115e-02,	3.55825269e-01,
1.61213354e-01,	9.69685794e-02,	2.32264965e-01,	-6.93181380e-02,
1.86467739e-01,	2.03098589e-01,	-1.10349710e-02,	2.63095846e-01,
-2.48406147e-01,	-3.25418955e-02,	1.74487006e-01,	2.62780720e-02,
1.91428194e-01,	2.03493779e-01,	-8.84696022e-02,	3.35631012e-01,
6.29476544e-02,	2.28818932e-01,	-7.72255471e-02,	3.05353195e-01,
3.63634109e-02,	2.30128273e-01,	3.03522210e-01,	1.05913376e-01,
1.26693452e-02,	9.535111494e-02,	4.52766233e-01,	-4.37470263e-02,
3.05687630e-01,	3.57706117e-02,	1.82867743e-01,	2.10106289e-01,
-1.71378996e-01,	2.60157245e-01,	-3.38655420e-01,	3.36603939e-01,
-7.65297319e-02,	2.15165094e-01,	3.72947326e-02,	1.96608549e-01,
1.07172893e-01,	3.07687901e-01,	3.97760529e-01,	1.06797074e-03,
8.12866229e-02,	2.95445495e-01,	5.47994817e-02,	1.13818287e-01,
4.07117263e-01,	1.48860323e-01,	3.88471838e-02,	3.79029267e-02,
1.09895981e-01,	-4.30946471e-02,	3.30298512e-01,	1.07254284e-01,
-1.13032643e-02,	-3.69192632e-02,	2.87732288e-01,	9.91961213e-02,
2.12225886e-01,	3.88660531e-01,	3.15623317e-01,	1.80996998e-01,
2.69970366e-01,	2.81850174e-01,	2.49972461e-01,	-2.33065542e-03,
2.34240860e-01,	1.51536128e-01,	6.56810225e-02,	1.35221573e-02,
3.03956323e-02,	9.22075626e-02,	1.28297232e-01,	2.04669352e-01,
2.26917512e-01,	-1.62627965e-01,	2.95984225e-01,	1.80934145e-01,
-6.34810776e-02,	4.36092057e-02,	1.39814157e-01,	1.72029014e-01,
1.65538329e-01,	2.24411690e-01,	2.15315070e-01,	1.16342630e-01,
-6.24745967e-021)			

y_test

```
from sklearn import metrics
print(metrics.r2_score(y_test,y_pred_rg))
print(metrics.r2_score(y_train,ridgecv.predict(x_train)))
```

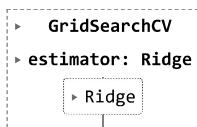
0.21073458438815917
 0.2061567210285109

▼ LASSO

```
from sklearn.linear_model import Lasso
from sklearn.model_selection import GridSearchCV

la=Ridge()

parametres={"alpha":[1,2,3,5,10,20,30,40,60,70,80,90]}
ridgecv=GridSearchCV(la,parametres,scoring="neg_mean_squared_error",cv=5)
ridgecv.fit(x_train,y_train)
```



```
print(ridgecv.best_params_)
```

```
{'alpha': 90}
```

```
print(ridgecv.best_score_)
```

```
-0.1139062113923418
```

```
y_pred_la=ridgecv.predict(x_test)
```

```
y_pred_la
```

```
array([-1.34413485e-01,  2.22561818e-01,  3.41692977e-01,  3.88209867e-03,
       4.84617338e-01,  1.16361483e-01,  3.30449743e-01,  1.27358807e-01,
      -1.34442619e-01,  3.77692888e-01,  1.33001445e-01,  2.69898751e-01,
     -2.54707392e-02,  5.25771894e-01,  2.67543514e-01,  2.78725024e-02,
      1.82233111e-01,  2.78896415e-01,  9.12689699e-02,  2.11494641e-01,
      2.70103341e-01,  8.44922044e-03,  8.74746722e-02,  1.05348798e-01,
     4.87749940e-01,  2.83080512e-01,  8.80556209e-02,  1.23817268e-01,
     4.82185624e-01,  9.34824523e-02,  -7.16448509e-02,  4.07003104e-02,
     1.08437994e-01,  3.42151399e-01,  1.22270929e-01,  6.85889862e-02,
    1.06690533e-01,  7.08689637e-02,  7.51570276e-02,  6.05829413e-02,
    1.08782897e-02,  -6.91368661e-03,  5.83191600e-02,  -1.54680056e-02,
   -4.02267475e-03,  4.08010612e-01,  3.43668700e-01,  -1.83519405e-01,
    5.29536511e-01,  4.27646098e-01,  1.95234877e-01,  4.25012930e-01,
    1.40754410e-01,  3.52173952e-01,  4.70372694e-01,  2.89240343e-01,
   -3.11642726e-02,  3.04206456e-01,  9.89337674e-03,  2.44569884e-01,
   -1.40249115e-02,  2.75133912e-01,  8.64669565e-02,  1.24214885e-01,
    3.48994545e-01,  3.41026778e-02,  3.40548051e-01,  1.95847356e-01,
    1.30040885e-01,  1.32259137e-01,  -2.34680143e-02,  3.04595468e-01,
    1.12452197e-01,  1.30525275e-01,  2.19329505e-01,  9.44722098e-02,
    9.98185782e-02,  2.60042486e-01,  2.51475715e-01,  4.59039018e-02,
   -7.94007856e-02,  -7.05812314e-03,  2.04344419e-01,  -3.97180151e-03,
   -5.91286905e-03,  1.26797761e-01,  8.02495203e-02,  2.55422079e-02,
    4.65384158e-02,  2.32985240e-01,  3.16063931e-01,  2.02833301e-01,
    3.14235904e-01,  2.33427181e-01,  -1.42446708e-02,  2.24789285e-01,
    2.94719863e-01,  2.94698323e-01,  1.19907108e-01,  9.47152062e-02,
    2.86026178e-01,  4.75925979e-01,  2.87802013e-01,  6.72561468e-03,
    1.65013565e-01,  1.72887026e-02,  3.34684186e-02,  2.15466121e-01,
    7.50317322e-02,  1.67646673e-01,  1.16585544e-01,  1.07157808e-01,
   -1.84689359e-02,  1.86217544e-01,  1.16586463e-01,  4.67201201e-02,
    1.11060472e-01,  2.27053971e-01,  -7.00247692e-02,  -5.81070776e-02,
    2.03141688e-01,  4.69029664e-02,  1.31525768e-01,  5.66738022e-01,
    2.41883060e-02,  -3.41250985e-02,  -1.13904557e-01,  2.18572744e-01,
    2.60568042e-01,  1.65533667e-01,  -5.94078459e-05,  2.60009384e-01,
    4.20709666e-01,  3.71031267e-01,  1.70250288e-01,  4.03052216e-01,
    4.67312765e-01,  1.98845366e-01,  1.55005619e-01,  3.41505080e-01,
    2.20024496e-01,  1.40989758e-01,  1.97796963e-01,  2.57841889e-01,
    2.99122317e-01,  9.24907038e-03,  1.39162817e-01,  -1.13916709e-01,
    1.97670909e-01,  2.70864780e-01,  1.22454317e-01,  2.58893294e-01,
    6.78818374e-02,  3.08485027e-01,  1.49347982e-01,  2.01436659e-02,
   -3.30262214e-02,  1.44305312e-01,  -8.99199978e-02,  -3.74712872e-02,
    3.10198738e-01,  -7.96862570e-02,  2.18579680e-01,  5.85363859e-01,
    1.98166099e-01,  3.02558934e-01,  1.82182301e-01,  1.84955080e-01,
    1.83694574e-02,  -7.41419216e-02,  4.48013268e-02,  1.38405390e-01,
    1.84013774e-01,  1.60373463e-01,  -6.83819091e-02,  2.00146771e-01,
    1.97563797e-01,  1.73505024e-01,  1.01481984e-01,  1.83169586e-01,
    1.99747065e-02,  1.81881922e-01,  -5.23948254e-02,  5.46171171e-01,
    6.66114639e-02,  5.88865384e-02,  3.17247692e-01,  9.77721299e-02,
    5.25297461e-01,  1.62566350e-01,  3.51341492e-01,  3.58324715e-01,
    2.37059552e-01,  8.05788438e-02,  1.36041888e-01,  2.66653277e-01,
    7.95513973e-02,  -6.95788172e-02,  3.29442074e-01,  8.93231393e-02,
    2.16673846e-01,  2.50725892e-01,  4.72995721e-01,  1.26285837e-01,
    2.72059331e-01,  6.13056795e-02,  4.38912502e-01,  -7.79381284e-02,
    2.09974643e-01,  2.20746796e-02,  1.56186553e-01,  2.26485767e-01,
    9.61150570e-02,  1.27870464e-01,  2.13995902e-01,  -9.95070059e-03,
    2.59908614e-02,  1.24499158e-01,  3.31256404e-02,  2.39369272e-01,
    1.48870840e-01,  2.49438253e-01,  5.25239856e-01,  1.25104891e-01,
    3.65711314e-01,  -5.96554519e-02,  1.45443911e-01,  2.80327834e-01,
    3.15149040e-01,  2.25038913e-02,  -2.55968584e-02,  3.39893959e-01,
```

```
from sklearn import metrics
print(metrics.r2_score(y_test,y_pred_la))
print(metrics.r2_score(y_train,ridgecv.predict(x_train)))
```

```
0.21073458438815917
0.2061567210285109
```

DECISION TREE

```
from sklearn.tree import DecisionTreeClassifier
dtc=DecisionTreeClassifier()
```



```
col_0    0   1

print(classification_report(y_test,pred))

      precision    recall  f1-score   support

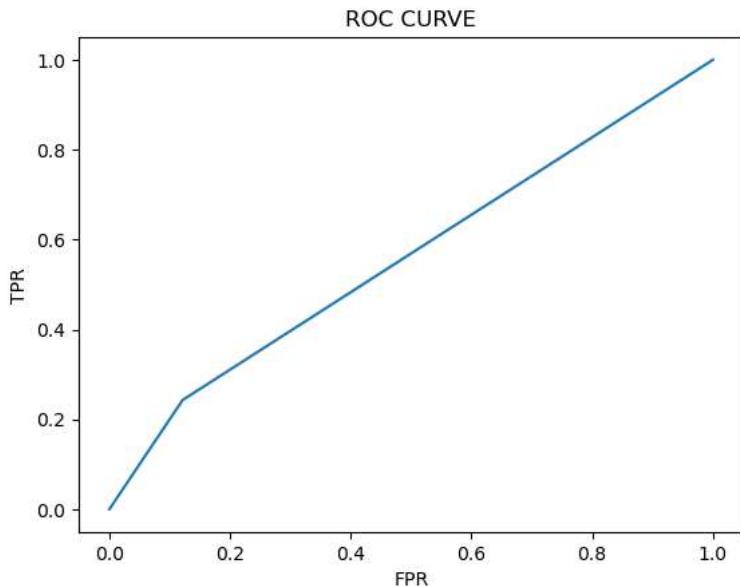
          0       0.86      0.88      0.87      371
          1       0.27      0.24      0.26       70

   accuracy                           0.78      441
    macro avg       0.57      0.56      0.56      441
weighted avg       0.77      0.78      0.77      441
```

```
probability=dtc.predict_proba(x_test)[:,1]
```

```
# roc_curve
fpr,tpr,thresholds = roc_curve(y_test,probability)

plt.plot(fpr,tpr)
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('ROC CURVE')
plt.show()
```



▼ RANDOM FOREST

```
from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()

forest_params = [ {'max_depth': list(range(10, 15)), 'max_features': list(range(0,14))}]

from sklearn.model_selection import GridSearchCV

rfc_cv= GridSearchCV(rfc,param_grid=forest_params,cv=10,scoring="accuracy")

rfc_cv.fit(x_train,y_train)
```



```
C:\ProgramData\anaconda3\Lib\site-packages\sklearn\model_selection\_validation.py:425: FitFailedWarning: 50 fits failed out of a total of 700.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.
```

Below are more details about the failures:

```
50 fits failed with the following error:
Traceback (most recent call last):
  File "C:\ProgramData\anaconda3\Lib\site-packages\sklearn\model_selection\_validation.py", line 425, in fit
    estimator.fit(X_train, y_train, **fit_params)
  File "C:\ProgramData\anaconda3\Lib\site-packages\sklearn\base.py", line 1144, in wrap_fit_method
    estimator._validate_params()
  File "C:\ProgramData\anaconda3\Lib\site-packages\sklearn\base.py", line 637, in _validate_params
    validate_parameter_constraints()
  File "C:\ProgramData\anaconda3\Lib\site-packages\sklearn\utils\_param_validation.py", line 100, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of RandomForestClassifier is not valid for this estimator. It must be a float between 0.0 and 1.0 or an integer between 0 and n_features, but got 4.0.
```

warnings.warn(some_fits_failed_message, FitFailedWarning)

C:\ProgramData\anaconda3\Lib\site-packages\sklearn\model_selection_search.py:976: UserWarning: 0.85323625 0.85226537 0.84934323 0.8512945 0.85324576 0.84837236
0.85423567 0.85713878 nan 0.84159528 0.84644965 0.85227489
0.85226537 0.85422616 0.85228441 0.85227489 0.85226537 0.85130402
0.84256615 0.85224634 0.84642109 0.84838188 nan 0.84936227
0.8483914 0.847411 0.85322673 0.85519703 0.85324576 0.85421664
0.847411 0.85615839 0.84935275 0.84545974 0.84642109 0.8561679
nan 0.84645917 0.85032362 0.85228441 0.85033314 0.85422616
0.85422616 0.85323625 0.85226537 0.84935275 0.84837236 0.8512945
0.8541976 0.847411 nan 0.84255663 0.84644013 0.85226537
0.85810965 0.85324576 0.8512945 0.84838188 0.85130402 0.8512945
0.85227489 0.8512945 0.84837236 0.84935275]

```
pred=rfc_cv.predict(x_test)
```

```
! ! !
```

```
print(classification_report(y_test,pred))
```

	precision	recall	f1-score	support
0	0.86	0.99	0.92	371
1	0.73	0.16	0.26	70
accuracy			0.86	441
macro avg	0.80	0.57	0.59	441
weighted avg	0.84	0.86	0.82	441

```
rfc_cv.best_params_
```

```
{'max_depth': 14, 'max_features': 4}
```

```
rfc_cv.best_score_
```

```
0.8581096516276412
```